

Abdul Rahiman Mihad

AI Engineer · Real-Time Conversational AI · Agentic Systems · MLOps · Abu Dhabi, UAE
+971 50 670 8992 | abdurahimanmihad@gmail.com | [Portfolio](#) | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

AI Engineer with production experience across two UAE and international technology companies, specialising in real-time conversational AI, agentic LLM systems, and full-stack AI product development. Built a voice-based AI interviewer with sub-second latency at HireUp, and currently architect the AI and backend for Vendoor, a B2B supplier platform. Skilled in Machine Learning, Generative AI, Large Language Models, NLP, Computer Vision, and MLOps, with cloud experience across AWS, Microsoft Azure, and Google Cloud. Microsoft Azure certified. Based in Abu Dhabi, available immediately.

CORE COMPETENCIES

Artificial Intelligence (AI) · Machine Learning (ML) · Deep Learning · Generative AI · Large Language Models (LLM) · Natural Language Processing (NLP) · Computer Vision · Conversational AI · Voice AI · AI Agents · Agentic AI · RAG · Prompt Engineering · LLM Fine-Tuning · Full-Stack Development · API Development · MLOps · Model Deployment · Data Science · Cloud Computing · Agile

WORK EXPERIENCE

Full-Stack AI Engineer, Vendoor

Jul 2026 to Present

B2B supplier discovery and tender platform for the UAE food and beverage market · Abu Dhabi, UAE (startup)

- Built the core **AI extraction engine** using Claude with forced tool-use for strict JSON output, automatically converting buyer documents and supplier price lists into structured tender line items and product catalogs.
- Engineered an **automated WhatsApp and email supplier workflow**: inbound message ingestion, AI-powered quote parsing, and confidence-scored matching that turns unstructured supplier replies into comparable line-item quotes.
- Designed the **full backend and data layer** in Node.js, Express, and MySQL, including buyer authentication, a founder admin console, and a complete tender lifecycle from request to award.
- Implemented **production reliability safeguards**: atomic processing locks to prevent duplicate quotes, graceful degradation when services are unavailable, and automated recovery of stalled messages.

AI/ML Engineer, HireUp

May 2026 to Jul 2026

Technology company headquartered in Japan (remote role) · app.hireup.now · Built “Alpha”, a real-time AI Voice Interviewer

- Designed the **full-duplex real-time voice pipeline**: WebSocket streaming with concurrent async loops, Deepgram Flux streaming speech-to-text, WebRTC voice detection, and streaming text-to-speech; **cut startup latency by about 1.5 seconds**.
- **Improved interview navigation accuracy from about 80% to 100%** by replacing unreliable LLM logic with a deterministic resolver, verified across live runs and automated tests.
- **Cut AI operating costs significantly** by redesigning how the system uses Large Language Models, removing unnecessary processing from every conversation turn.
- Built an **automated evaluation suite** that replays full interviews in seconds, plus MLflow and LangSmith monitoring for tracking quality across software updates.

KEY PROJECTS

Self-Corrective RAG System

 LangGraph · AWS Bedrock · FastAPI · MLflow · Docker [Live Demo](#)

- **Reduced hallucinations 40%** and reached **92% query accuracy** with a self-correcting multi-agent pipeline: document grading, hallucination detection, and query rewriting using Llama 3.1 (70B) on AWS Bedrock.

End-to-End Litter Detection System

 YOLOv8 · Computer Vision · DVC · MLflow · Streamlit [Live Demo](#)

- **Boosted small-object detection accuracy 15%** on **1,500+ hand-labelled images**, tracked across 50+ experiments with DVC and MLflow, and deployed as a live web application.

Language Model Built from Scratch Deep Learning · PyTorch

- Built a small language model from scratch, including a **custom automatic differentiation engine**, to master how systems like ChatGPT work internally.

EDUCATION

Bachelor of Computer Applications (BCA) Aug 2022 to Nov 2025
Yenepoya University, India · Specialisation: Machine Learning, AI and Robotics · Industry partnership with TCS and Microsoft

CERTIFICATIONS

- **Microsoft Azure:** AI Fundamentals (AI-900) · Azure Fundamentals (AZ-900) · Data Fundamentals (DP-900)
- **IBM / Coursera, Generative AI:** [Advanced LLM Fine-Tuning](#) · [Fine-Tuning Transformers](#) · [AI Agents: RAG and LangChain](#) · [LLM Architecture and Data Prep](#) · [LM with Transformers](#) · [Foundational Models for NLP](#) · [EDA for ML](#)
- **Kaggle:** Advanced SQL · Intro to SQL

SKILLS

AI / LLM	LLMs · RAG · Multi-Agent Systems · LangGraph · LangChain · Hugging Face · Prompt Engineering · PEFT/LoRA · Transformers · NLP · Computer Vision
Real-Time AI	WebSockets · asyncio · Deepgram Speech-to-Text and Text-to-Speech · WebRTC · Streaming Pipelines
Engineering	Python · JavaScript · Node.js · SQL · PyTorch · TensorFlow · FastAPI · Express · REST APIs · Streamlit
MLOps	MLflow · LangSmith · DVC · Docker · GitHub Actions · CI/CD · Model Deployment · Model Monitoring
Cloud / Data	AWS (Bedrock, EC2, S3) · Microsoft Azure · Google Cloud · MySQL · PostgreSQL · Redis · BigQuery
Languages:	English (Fluent) · Malayalam (Native) · Hindi (Elementary) · Arabic (Elementary)